

Final Project 1361

Daniel Antantis

```
library(tidyverse)
library(glmnet)
library(gam)
library(boot)
library(leaps)
library(randomForest)
library(caret)
```

Read in Data

```
raw_data = read_csv("~/Documents/STAT1361/train.csv")

## Rows: 6552 Columns: 15

## -- Column specification -----
## Delimiter: ","
## chr (4): Date, Seasons, Holiday, Functioning
## dbl (11): Count, Hour, Temperature, Humidity, Wind, Visibility, Dew, Solar, ...

##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

Data Wrangling

```
#Separate the date variable into three separate variables because I want to see if individually they are
out.data = raw_data %>%
  separate(Date, sep="/", into = c("Day", "Month", "Year"))
```

```
summary(out.data$Count)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      0.0   189.0   492.0   702.9  1062.0  3556.0
```

```
#Create percentiles to extract outliers in the data with extreme count values
lower_bound = quantile(out.data$Count, 0.01)
upper_bound = quantile(out.data$Count, 0.99)

outliers.ind = which(out.data$Count < lower_bound | out.data$Count > upper_bound)

out.data[outliers.ind, ]
```

```
## # A tibble: 66 x 17
##   Count Day   Month Year   Hour Temperature Humidity Wind Visibility Dew
##   <dbl> <chr> <chr> <chr> <dbl>      <dbl>    <dbl> <dbl>    <dbl> <dbl>
## 1  2692 16     4     2018    18        17      28    3.1     2000  -1.6
## 2  2807 25     4     2018    18       21.2     32    3.8     1927   3.8
## 3  2574 26     4     2018    18       17.4     45    3.1     1092   5.3
## 4  2661 4      5     2018    18       17.1     35    3.4     1961   1.4
## 5  3130 9      5     2018    18       20.6     41    2.3     2000   6.8
## 6  2701 11     5     2018    18       17.9     37    3.1     1819   2.9
## 7  2906 14     5     2018    18       23.6     48    3.1      666  11.9
## 8  3069 21     5     2018    18       21.6     48    2.5     1884  10.1
## 9  3123 23     5     2018    18       21.7     40    3.5     1987   7.4
## 10 2916 25     5     2018    18       23.3     32    2.6     1772   5.6
## # ... with 56 more rows, and 7 more variables: Solar <dbl>, Rainfall <dbl>,
## #   Snowfall <dbl>, Seasons <chr>, Holiday <chr>, Functioning <chr>, ID <dbl>
```

```
#Explore the outliers and look for trends that may help explain the extreme values
range(out.data[outliers.ind, ]$Temperature)
```

```
## [1] 14.3 33.0
```

```
range(out.data[outliers.ind, ]$Humidity)
```

```
## [1] 27 77
```

```
range(out.data[outliers.ind, ]$Hour)
```

```
## [1] 17 20
```

```
mean(out.data[outliers.ind, ]$Temperature)
```

```
## [1] 24.97121
```


```
mean(out.data[outliers.ind, ]$Humidity)
```

```
## [1] 49.18182
```

```
#Remove the outliers from the dataset
data = out.data[-outliers.ind, ]
```

Colinearity


```
#Explore the strong colinearity between dew point temperature and temperature as they are both measuring
plot(data$Dew, data$Temperature)
```




Relationship Analysis

Relationship Analysis

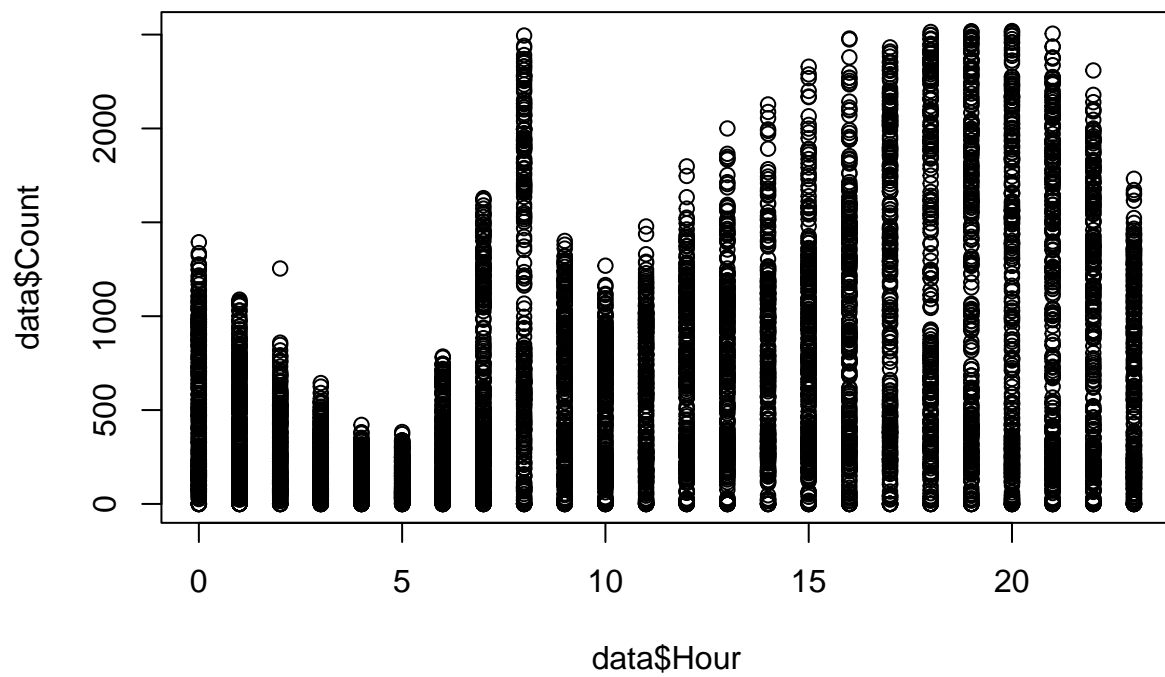
```
plot(data$Temperature, data$Count)
```



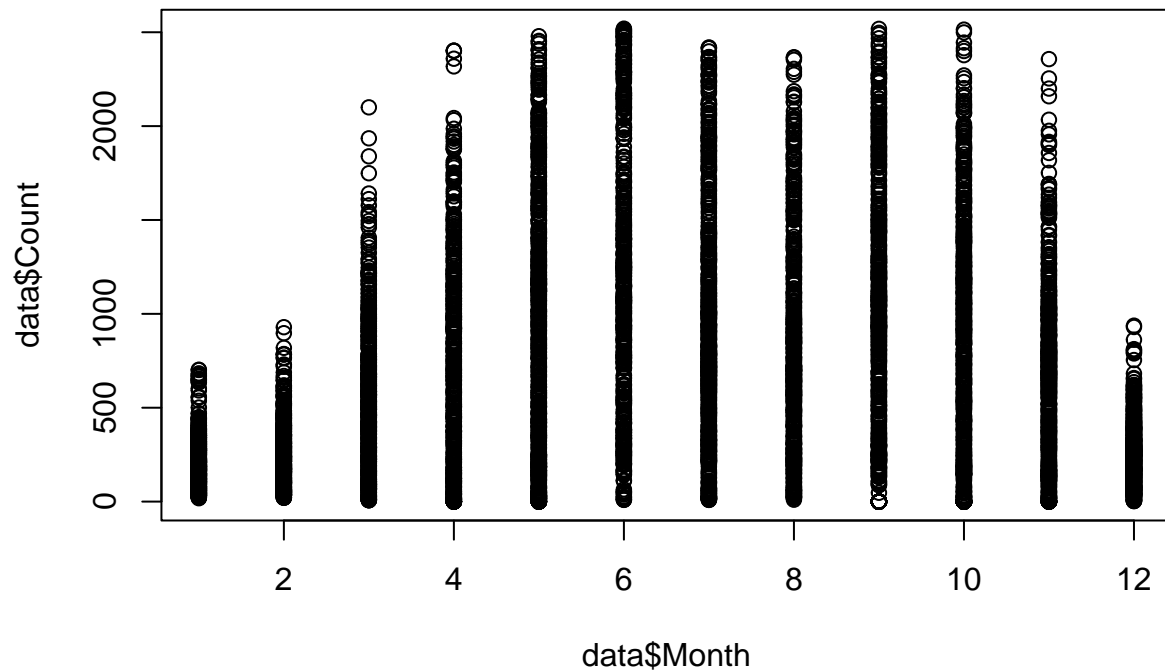
```
plot(data$Rainfall, data$Count)
```




```
plot(data$Hour, data$Count)
```




```
plot(data$Month, data$Count)
```




#Explore the functioning variable and see that it has a very strong affect on bike rental count therfor
`plot(as.factor(data$Functioning), data$Count)`



#Extra plots of some of the categorical variables
`plot(as.factor(data$Holiday), data$Count)`



```
plot(as.factor(data$Seasons), data$Count)
```



Data

Split

```
#Splitting the data into training and testing dataset with the training set containing 75% of the data
set.seed(21)
sampleSize = floor(0.75 * nrow(data))
split = sample(seq_len(nrow(data)), size = sampleSize)
train = data[split, ]
test = data[-split, ]
```

Linear Model

#Test all of the variables in the dataset and see which ones standout

```
lin.fit.test = lm(Count~Hour+Temperature+Humidity+Wind+Visibility+Dew+Solar+Rainfall+Snowfall+Holiday+Month+Year+Day+Functioning, data = train)
summary(lin.fit.test)
```

```
##
## Call:
## lm(formula = Count ~ Hour + Temperature + Humidity + Wind + Visibility +
##     Dew + Solar + Rainfall + Snowfall + Holiday + Month + Year +
##     Day + Functioning, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1264.40  -251.12   -39.11   207.74  1561.27
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -760.55654   118.80561   -6.402 1.68e-10 ***
## Hour           24.96820     0.90374   27.628 < 2e-16 ***
## Temperature    25.55621     4.39474    5.815 6.45e-09 ***
## Humidity       -7.49874     1.20431   -6.227 5.17e-10 ***
## Wind           12.03098     6.04632    1.990 0.046669 *
## Visibility      0.07947     0.01456    5.458 5.04e-08 ***
## Dew            4.12105     4.58771    0.898 0.369081
## Solar          -70.07145     9.24739   -7.577 4.20e-14 ***
## Rainfall       -47.40677     4.86234   -9.750 < 2e-16 ***
## Snowfall        47.59727    12.76705    3.728 0.000195 ***
## HolidayNo Holiday  83.51357    28.59155    2.921 0.003506 **
## Month10         332.23645    39.77814    8.352 < 2e-16 ***
## Month11         257.20590    32.35118    7.950 2.30e-15 ***
## Month12          47.15079    26.76941    1.761 0.078240 .
## Month2          -31.31106    27.99987   -1.118 0.263513
## Month3          103.44028    32.95394    3.139 0.001706 **
## Month4          200.71836    38.00621    5.281 1.34e-07 ***
## Month5          285.85742    42.40298    6.741 1.75e-11 ***
## Month6          355.34173    49.57794    7.167 8.81e-13 ***
## Month7           4.18473    56.55410    0.074 0.941017
## Month8         -167.12984    59.31014   -2.818 0.004854 **
## Month9          239.27465    49.95153    4.790 1.72e-06 ***
## Year2018                NA                NA                NA                NA
## Day10             14.82557    42.80940    0.346 0.729121
## Day11             98.76443    47.25837    2.090 0.036681 *
## Day12            -17.47042    44.86352   -0.389 0.696988
## Day13            144.16322    41.48588    3.475 0.000515 ***
## Day14             85.66177    42.42476    2.019 0.043527 *
## Day15            -5.47481    42.54965   -0.129 0.897625
## Day16             22.80189    43.24139    0.527 0.597998
## Day17             13.28926    42.81167    0.310 0.756261
## Day18            -29.98797    42.19753   -0.711 0.477331
## Day19             64.45597    42.60567    1.513 0.130384
## Day2             -61.56587    45.12660   -1.364 0.172539
## Day20             14.31648    41.01831    0.349 0.727085
## Day21            -14.56142    44.20499   -0.329 0.741863
```

```
## Day22          -8.87794  46.64303 -0.190 0.849052
## Day23          -15.70737  43.35961 -0.362 0.717175
## Day24          -92.18446  43.67347 -2.111 0.034844 *
## Day25           34.52368  40.15148  0.860 0.389922
## Day26          -83.67984  43.15031 -1.939 0.052528 .
## Day27           13.62328  42.39415  0.321 0.747961
## Day28          -75.13244  42.09606 -1.785 0.074359 .
## Day29          -32.81023  44.27366 -0.741 0.458682
## Day3           113.47540  48.89570  2.321 0.020341 *
## Day30           11.48323  50.68646  0.227 0.820780
## Day31          -31.70084  53.29614 -0.595 0.552002
## Day4           -0.95564  40.80279 -0.023 0.981315
## Day5           18.04469  42.67138  0.423 0.672405
## Day6           29.07477  40.95955  0.710 0.477837
## Day7           105.63816  45.28931  2.333 0.019714 *
## Day8           15.69984  46.30124  0.339 0.734564
## Day9           40.59825  41.30816  0.983 0.325748
## FunctioningYes  974.15560  31.99808 30.444 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 386 on 4811 degrees of freedom
## Multiple R-squared:  0.607, Adjusted R-squared:  0.6027
## F-statistic: 142.9 on 52 and 4811 DF, p-value: < 2.2e-16
```

#The final linear model chosen to represent the data

```
lin.fit = lm(Count~Hour+Temperature+Humidity+Visibility+Solar+Rainfall+Snowfall+Holiday+Month+Day+FunctioningYes, data=train)
summary(lin.fit)
```

```
##
## Call:
## lm(formula = Count ~ Hour + Temperature + Humidity + Visibility +
##      Solar + Rainfall + Snowfall + Holiday + Month + Day + Functioning,
##      data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1262.46  -251.26   -38.17   209.38  1553.56
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -827.20366    67.96109  -12.172 < 2e-16 ***
## Hour             25.38570     0.87557   28.993 < 2e-16 ***
## Temperature     29.18442     1.49524   19.518 < 2e-16 ***
## Humidity        -6.55041     0.49361  -13.270 < 2e-16 ***
## Visibility       0.08148     0.01451    5.617 2.05e-08 ***
## Solar          -67.75729     8.76779   -7.728 1.32e-14 ***
## Rainfall       -47.48225     4.83397   -9.823 < 2e-16 ***
## Snowfall        47.11824    12.75186    3.695 0.000222 ***
## HolidayNo Holiday  82.39427    28.59451    2.881 0.003976 **
## Month10          330.18588    39.63599    8.330 < 2e-16 ***
## Month11          253.13946    32.29785    7.838 5.61e-15 ***
## Month12           45.25413    26.76057    1.691 0.090887 .
## Month2          -31.85186    27.98843   -1.138 0.255162
```



```

## Month3      106.11039   32.93354   3.222 0.001282 **
## Month4      202.83947   37.99601   5.338 9.81e-08 ***
## Month5      287.57442   42.17512   6.819 1.03e-11 ***
## Month6      358.39184   49.33155   7.265 4.33e-13 ***
## Month7        6.57068   56.05174   0.117 0.906687
## Month8     -162.43348   58.97468  -2.754 0.005904 **
## Month9      238.15703   49.65426   4.796 1.66e-06 ***
## Day10        17.82049   42.79883   0.416 0.677152
## Day11       103.07326   47.20863   2.183 0.029058 *
## Day12       -21.20319   44.83229  -0.473 0.636274
## Day13       142.53796   41.47330   3.437 0.000593 ***
## Day14        84.01504   42.40193   1.981 0.047604 *
## Day15       -8.40786   42.53651  -0.198 0.843318
## Day16        20.81610   43.23398   0.481 0.630201
## Day17        10.38756   42.80054   0.243 0.808250
## Day18       -37.96349   42.04372  -0.903 0.366596
## Day19        61.41300   42.58258   1.442 0.149308
## Day2       -65.37097   45.10483  -1.449 0.147316
## Day20        13.49517   40.99847   0.329 0.742047
## Day21       -19.28412   44.13812  -0.437 0.662201
## Day22        -8.33717   46.64414  -0.179 0.858149
## Day23       -14.20174   43.34239  -0.328 0.743180
## Day24       -90.89819   43.65635  -2.082 0.037383 *
## Day25        32.95142   40.15176   0.821 0.411874
## Day26       -82.95918   43.15008  -1.923 0.054593 .
## Day27        11.85714   42.39659   0.280 0.779741
## Day28       -77.93300   42.07997  -1.852 0.064084 .
## Day29       -34.42829   44.26885  -0.778 0.436779
## Day3        111.05493   48.87439   2.272 0.023115 *
## Day30         8.51746   50.68139   0.168 0.866544
## Day31       -32.82581   53.30368  -0.616 0.538038
## Day4         0.37183   40.80954   0.009 0.992731
## Day5        21.23445   42.65761   0.498 0.618656
## Day6        30.02392   40.96854   0.733 0.463684
## Day7       107.97817   45.28481   2.384 0.017144 *
## Day8        14.84240   46.30875   0.321 0.748596
## Day9        36.83117   41.28071   0.892 0.372324
## FunctioningYes 969.54352  31.93579  30.359 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 386.1 on 4813 degrees of freedom
## Multiple R-squared:  0.6066, Adjusted R-squared:  0.6025
## F-statistic: 148.4 on 50 and 4813 DF, p-value: < 2.2e-16

```

```

lin.pred = predict(lin.fit, test, type = 'response')
mean((lin.pred - test$Count)^2)

```

```
## [1] 155040.3
```

Ridge

```

#In order to do ridge regression and lasso we must first convert the data into matrices and vectors
train.x = model.matrix(Count~Hour+Temperature+Humidity+Visibility+Solar+Rainfall+Snowfall+Holiday+Month+Day)
train.count = train$Count
test.x = model.matrix(Count~Hour+Temperature+Humidity+Visibility+Solar+Rainfall+Snowfall+Holiday+Month+Day)
test.count = test$Count

ridge.fit = cv.glmnet(train.x, train.count, alpha = 0)
ridge.lambda = ridge.fit$lambda.min

ridge.pred = predict(ridge.fit, s = ridge.lambda, newx = test.x)
mean((ridge.pred - test.count)^2)

```

```
## [1] 156542.1
```

Coefficients

```

#Examine the coefficient values of ridge regression and see where the model is weighted most heavily
ridge.coef = predict(ridge.fit, type = "coefficients", s = ridge.lambda)
ridge.coef

```

```

## 52 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept)  -756.41984357
## (Intercept)      .
## Hour          25.20441115
## Temperature   24.53450266
## Humidity      -5.29664146
## Visibility     0.08594388
## Solar         -35.78087928
## Rainfall      -46.65097937
## Snowfall      25.88587586
## HolidayNo Holiday 73.89189041
## Month10       291.57099432
## Month11       214.94656223
## Month12       -14.79950985
## Month2        -86.69024081
## Month3         65.21893162
## Month4        169.54792121
## Month5        270.30546867
## Month6        360.83456055
## Month7         38.73290977
## Month8       -113.91267232
## Month9        234.81899205
## Day10          3.29180952
## Day11          68.68700150
## Day12         -38.11303694
## Day13         120.04840758
## Day14          74.66014605
## Day15         -9.36159380
## Day16          13.70178452
## Day17           2.88485268
## Day18        -36.66302242

```

```
## Day19          51.59497727
## Day2          -82.69391669
## Day20          11.62590134
## Day21         -14.45294922
## Day22         -18.89173925
## Day23         -23.48640811
## Day24         -97.40460716
## Day25          24.72808813
## Day26         -81.62528099
## Day27           4.86088264
## Day28         -86.26413647
## Day29         -40.06885031
## Day3           92.95100726
## Day30          -0.85755125
## Day31         -26.51280702
## Day4          -17.40438166
## Day5           7.06520903
## Day6           7.73930999
## Day7           96.87295918
## Day8           -6.48585885
## Day9           11.45451269
## FunctioningYes 895.37553836
```

Lasso

```
lasso.fit = cv.glmnet(train.x, train.count, alpha = 1)
lasso.lambda = lasso.fit$lambda.min

lasso.pred = predict(lasso.fit, s = lasso.lambda, newx = test.x)
mean((lasso.pred - test.count)^2)
```

```
## [1] 155180.7
```

Coefficients

```
lasso.coef = predict(lasso.fit, type = "coefficients", s = lasso.lambda)
lasso.coef
```

```
## 52 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept) -798.7363757
## (Intercept) .
## Hour        25.1942493
## Temperature 29.8215795
## Humidity    -6.1829193
## Visibility   0.0848889
## Solar       -62.5334573
## Rainfall    -47.1312742
## Snowfall     39.6038593
## HolidayNo Holiday 76.7807117
## Month10      284.3979813
## Month11      218.8138534
## Month12       17.9091067
```

```
## Month2          -51.6216511
## Month3           67.0238220
## Month4          157.7576781
## Month5          239.7138166
## Month6          306.5343998
## Month7          -43.0377142
## Month8         -211.5701176
## Month9          186.2000560
## Day10            2.5363093
## Day11           87.8969593
## Day12          -23.2194963
## Day13          125.2025118
## Day14           65.7520486
## Day15          -12.8769955
## Day16           4.8249558
## Day17            .
## Day18          -37.5617098
## Day19           45.0951113
## Day2           -71.6810993
## Day20            .
## Day21          -20.9653362
## Day22          -11.9218835
## Day23          -15.8115625
## Day24          -92.3942631
## Day25           19.8928737
## Day26          -84.8017992
## Day27            .
## Day28          -81.4116723
## Day29          -37.0016254
## Day3            89.5340231
## Day30            .
## Day31          -35.2859332
## Day4            -6.2296599
## Day5             3.3947932
## Day6            13.0423591
## Day7            94.8612924
## Day8            .
## Day9            19.7400685
## FunctioningYes  959.3879449
```

Relaxed Lasso

```
relax.fit = glmnet(train.x, train.count, relax = TRUE)
relax.lambda = relax.fit$lambda.min
relax.pred = predict(relax.fit, test.x, s = relax.lambda)
mean((relax.pred - test.count)^2)
```

```
## [1] 189335.3
```

Non-Linear Methods

#Examine some of the variables in differing degrees of non-linear polynomial functions

```
fit = lm(Count~ poly(Hour,5)+poly(Temperature,5)+poly(Humidity,5)+poly(Visibility,5)+poly(Solar,5)+poly(Rainfall,5)+poly(Holiday)+poly(Month)+poly(Day)+poly(Functioning), data = train)
summary(fit)
```

```
##
## Call:
## lm(formula = Count ~ poly(Hour, 5) + poly(Temperature, 5) + poly(Humidity,
##      5) + poly(Visibility, 5) + poly(Solar, 5) + poly(Rainfall,
##      5) + Holiday + Month + Day + Functioning, data = train)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-1373.84	-206.48	-23.33	167.90	1412.08

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-522.484	50.169	-10.414	< 2e-16	***
poly(Hour, 5)1	10743.926	406.663	26.420	< 2e-16	***
poly(Hour, 5)2	7760.263	753.264	10.302	< 2e-16	***
poly(Hour, 5)3	-2116.244	441.944	-4.788	1.73e-06	***
poly(Hour, 5)4	-4273.820	464.609	-9.199	< 2e-16	***
poly(Hour, 5)5	-8012.762	372.212	-21.527	< 2e-16	***
poly(Temperature, 5)1	24045.224	1203.382	19.981	< 2e-16	***
poly(Temperature, 5)2	5348.191	559.158	9.565	< 2e-16	***
poly(Temperature, 5)3	-5061.365	407.517	-12.420	< 2e-16	***
poly(Temperature, 5)4	-4161.305	391.866	-10.619	< 2e-16	***
poly(Temperature, 5)5	-1349.479	356.559	-3.785	0.000156	***
poly(Humidity, 5)1	-6638.939	647.386	-10.255	< 2e-16	***
poly(Humidity, 5)2	-5358.995	427.071	-12.548	< 2e-16	***
poly(Humidity, 5)3	-1927.284	387.663	-4.972	6.87e-07	***
poly(Humidity, 5)4	-2761.400	361.098	-7.647	2.47e-14	***
poly(Humidity, 5)5	202.253	344.120	0.588	0.556734	
poly(Visibility, 5)1	567.218	539.273	1.052	0.292935	
poly(Visibility, 5)2	750.053	375.508	1.997	0.045835	*
poly(Visibility, 5)3	-1418.907	349.690	-4.058	5.04e-05	***
poly(Visibility, 5)4	1401.978	343.362	4.083	4.52e-05	***
poly(Visibility, 5)5	-43.491	337.654	-0.129	0.897519	
poly(Solar, 5)1	5821.259	911.567	6.386	1.86e-10	***
poly(Solar, 5)2	-5953.504	459.049	-12.969	< 2e-16	***
poly(Solar, 5)3	4777.727	388.353	12.303	< 2e-16	***
poly(Solar, 5)4	-3120.439	357.177	-8.736	< 2e-16	***
poly(Solar, 5)5	2618.301	345.953	7.568	4.51e-14	***
poly(Rainfall, 5)1	-2441.856	391.958	-6.230	5.07e-10	***
poly(Rainfall, 5)2	1889.200	371.366	5.087	3.77e-07	***
poly(Rainfall, 5)3	-1608.831	354.828	-4.534	5.93e-06	***
poly(Rainfall, 5)4	2118.304	345.385	6.133	9.31e-10	***
poly(Rainfall, 5)5	-1511.674	335.326	-4.508	6.70e-06	***
HolidayNo Holiday	56.287	24.342	2.312	0.020803	*
Month10	472.492	37.022	12.763	< 2e-16	***
Month11	443.926	29.873	14.860	< 2e-16	***
Month12	126.993	23.465	5.412	6.53e-08	***
Month2	33.499	24.508	1.367	0.171745	
Month3	211.680	30.533	6.933	4.67e-12	***

```

## Month4          354.122    35.841    9.880 < 2e-16 ***
## Month5          350.522    40.060    8.750 < 2e-16 ***
## Month6          251.779    45.860    5.490 4.22e-08 ***
## Month7          -74.338    52.062   -1.428 0.153393
## Month8         -174.662    54.428   -3.209 0.001341 **
## Month9          221.587    45.730    4.846 1.30e-06 ***
## Day10           5.380     36.520    0.147 0.882893
## Day11          65.248     40.303    1.619 0.105527
## Day12          12.346     38.248    0.323 0.746871
## Day13          88.273     35.312    2.500 0.012459 *
## Day14          18.713     36.138    0.518 0.604601
## Day15         -30.586     36.214   -0.845 0.398385
## Day16          12.170     36.979    0.329 0.742080
## Day17         -13.684     36.486   -0.375 0.707633
## Day18         -22.964     35.758   -0.642 0.520764
## Day19          59.551     36.505    1.631 0.102887
## Day2          -60.456     38.596   -1.566 0.117327
## Day20          26.274     35.008    0.751 0.452982
## Day21          24.875     37.766    0.659 0.510150
## Day22          35.277     39.963    0.883 0.377423
## Day23          -0.814     36.966   -0.022 0.982431
## Day24         -89.936     37.294   -2.412 0.015921 *
## Day25         -17.995     34.367   -0.524 0.600576
## Day26        -104.975     37.075   -2.831 0.004654 **
## Day27          20.499     36.192    0.566 0.571151
## Day28         -67.759     35.933   -1.886 0.059393 .
## Day29         -39.281     37.702   -1.042 0.297522
## Day3          126.519     41.792    3.027 0.002480 **
## Day30         -48.737     43.201   -1.128 0.259310
## Day31           4.024     44.932    0.090 0.928644
## Day4          -12.447     34.788   -0.358 0.720516
## Day5          -19.006     36.338   -0.523 0.600986
## Day6           10.004     34.961    0.286 0.774773
## Day7           76.523     38.700    1.977 0.048059 *
## Day8           52.969     39.572    1.339 0.180781
## Day9           36.000     35.293    1.020 0.307760
## FunctioningYes  994.926     27.316   36.422 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 327.4 on 4790 degrees of freedom
## Multiple R-squared:  0.7186, Adjusted R-squared:  0.7143
## F-statistic: 167.5 on 73 and 4790 DF, p-value: < 2.2e-16

```

#Compare the different functions on each variable and see which produces the most reliable model

```

anova1 = gam(Count~Hour+Temperature+Humidity+Visibility+Solar+Rainfall+Holiday+Month+Day+Functioning, data=train)
anova2 = gam(Count~ Hour+s(Temperature,3)+s(Humidity,2)+s(Visibility,2)+s(Solar,3)+Rainfall+Holiday+Month+Day+Functioning, data=train)
anova3 = gam(Count~ Hour+Temperature+s(Humidity,3)+s(Visibility,2)+s(Solar,3)+s(Rainfall,3)+Holiday+Month+Day+Functioning, data=train)
anova4 = gam(Count~ s(Hour,5)+s(Temperature,5)+s(Humidity,4)+s(Solar,4)+s(Rainfall,5)+Holiday+Month+Day+Functioning, data=train)
anova(anova1, anova2, anova3, anova4, test = 'F')

```

```
## Analysis of Deviance Table
```

```
##
```

```
## Model 1: Count ~ Hour + Temperature + Humidity + Visibility + Solar +
```

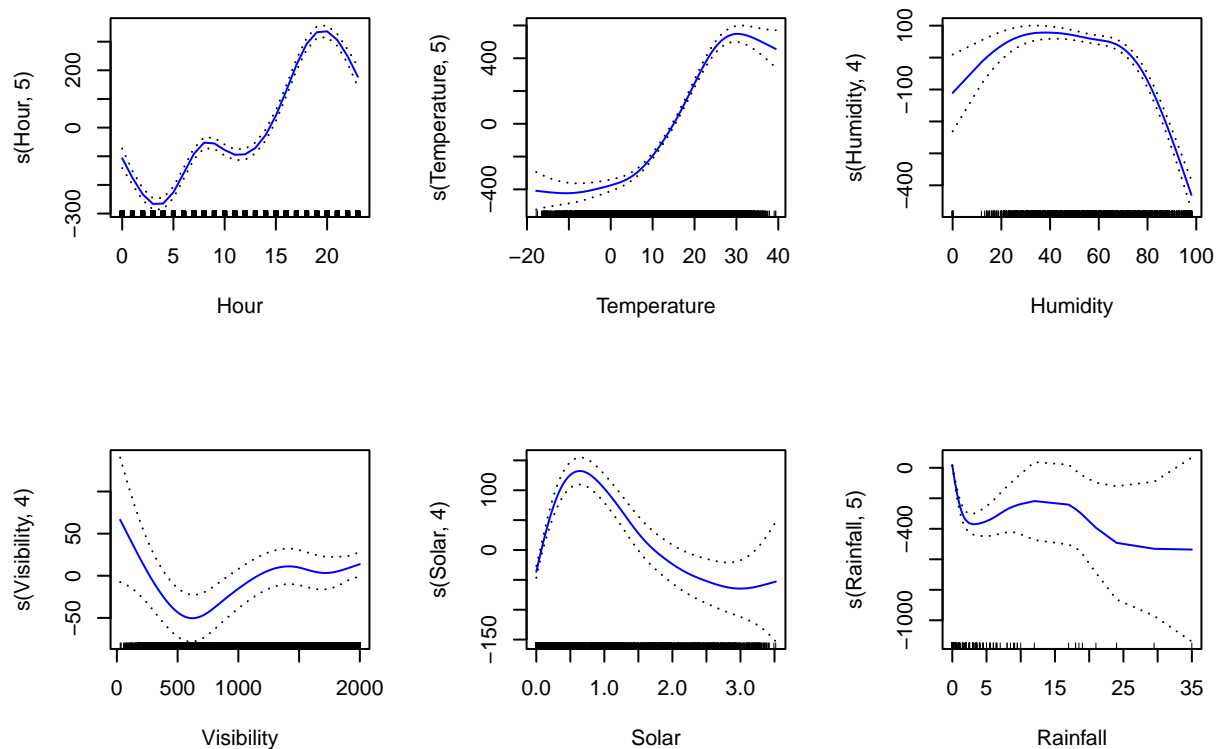
```
## Rainfall + Holiday + Month + Day + Functioning
## Model 2: Count ~ Hour + s(Temperature, 3) + s(Humidity, 2) + s(Visibility,
## 2) + s(Solar, 3) + Rainfall + Holiday + Month + Day + Functioning
## Model 3: Count ~ Hour + Temperature + s(Humidity, 3) + s(Visibility, 2) +
## s(Solar, 3) + s(Rainfall, 3) + Holiday + Month + Day + Functioning
## Model 4: Count ~ s(Hour, 5) + s(Temperature, 5) + s(Humidity, 4) + s(Solar,
## 4) + s(Rainfall, 5) + Holiday + Month + Day + Functioning
## Resid. Df Resid. Dev      Df Deviance      F      Pr(>F)
## 1      4814  719708622
## 2      4808  646305666 6.00023  73402956 111.60 < 2.2e-16 ***
## 3      4807  653791481 0.99986  -7485815
## 4      4797  525852566 9.99975 127938914 116.71 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
model = glm(Count~ s(Hour,5)+s(Temperature,5)+s(Humidity,4)+s(Solar,4)+s(Rainfall,5)+Holiday+Month+Day+
cv.glm(train, model, K = 10)$delta[1]
```

```
## [1] 152530.3
```

GAM

```
gam.mod = gam(Count~ s(Hour,5)+s(Temperature,5)+s(Humidity,4)+s(Visibility,4)+s(Solar,4)+s(Rainfall,5)+
par(mfrow = c(2,3))
plot(gam.mod, se = TRUE, col = 'blue')
```



```
summary(gam.mod)
```

```
##
## Call: gam(formula = Count ~ s(Hour, 5) + s(Temperature, 5) + s(Humidity,
##      4) + s(Visibility, 4) + s(Solar, 4) + s(Rainfall, 5) + Holiday +
##      Month + Day + Functioning, data = train)
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1296.59  -214.78   -33.08   160.09  1452.60
##
## (Dispersion Parameter for gaussian family taken to be 109182.6)
##
## Null Deviance: 1824319943 on 4863 degrees of freedom
## Residual Deviance: 523312098 on 4793 degrees of freedom
## AIC: 70302.09
##
## Number of Local Scoring Iterations: NA
##
## Anova for Parametric Effects
##           Df    Sum Sq   Mean Sq    F value    Pr(>F)
## s(Hour, 5)      1 280745339 280745339 2571.3384 < 2.2e-16 ***
## s(Temperature, 5) 1 455217965 455217965 4169.3280 < 2.2e-16 ***
## s(Humidity, 4)    1  50626536  50626536  463.6870 < 2.2e-16 ***
## s(Visibility, 4)  1   3178917   3178917   29.1156 7.147e-08 ***
## s(Solar, 4)       1    349792    349792    3.2037  0.07353 .
## s(Rainfall, 5)    1 10270516 10270516   94.0674 < 2.2e-16 ***
## Holiday          1  1880647   1880647   17.2248 3.378e-05 ***
## Month           11 125090856 11371896  104.1549 < 2.2e-16 ***
## Day             30 28385828   946194    8.6662 < 2.2e-16 ***
## Functioning      1 145549941 145549941 1333.0877 < 2.2e-16 ***
## Residuals       4793 523312098   109183
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Anova for Nonparametric Effects
##           Npar Df   Npar F      Pr(F)
## (Intercept)
## s(Hour, 5)      4 173.653 < 2.2e-16 ***
## s(Temperature, 5) 4  90.998 < 2.2e-16 ***
## s(Humidity, 4)   3  96.347 < 2.2e-16 ***
## s(Visibility, 4)  3   6.509 0.0002166 ***
## s(Solar, 4)      3  69.045 < 2.2e-16 ***
## s(Rainfall, 5)   4  35.804 < 2.2e-16 ***
## Holiday
## Month
## Day
## Functioning
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
#Output the R2 and the MSE
```

```
preds = predict(gam.mod, test)
RSS = sum((test$Count - preds)2)
```



```
TSS = sum((test$Count - mean(test$Count)) ^ 2)
1 - (RSS / TSS)
```

```
## [1] 0.6945151
```

```
mean((test$Count - preds)^2)
```

```
## [1] 115668.9
```

K-Nearest Neighbors

```
set.seed(21)
knn.model = knnreg(train.x, train.count)
knn.pred = predict(knn.model, data.frame(test.x))
mean((test$Count - knn.pred)^2)
```

```
## [1] 189226.8
```

Random Forests

Bagging

```
bag.fit = randomForest(Count~Hour+Temperature+Humidity+Visibility+Dew+Wind+Snowfall+Solar+Rainfall+Holiday)
bag.pred = predict(bag.fit, test)
mean((test$Count - bag.pred)^2)
```

```
## [1] 47658.63
```

```
#Output the variable importance in the bagged random forest model
importance(bag.fit)
```

```
##           %IncMSE IncNodePurity
## Hour           66.326960      489419808
## Temperature    46.138989      540185274
## Humidity        25.149905      176391217
## Visibility      14.514499      38300731
## Dew              7.780979      71410119
## Wind            7.281313      26918076
## Snowfall        5.132555      1082239
## Solar           13.176327      111204486
## Rainfall        14.534428      66212759
## Holiday         5.490983      3798820
## Month           5.671254      59630957
## Day             12.114436      33997108
## Functioning     55.342751      189910776
```

RF

```
rf.fit = randomForest(Count~Hour+Temperature+Humidity+Visibility+Dew+Wind+Snowfall+Solar+Rainfall+Holiday+Month+Day+Functioning)
rf.pred = predict(rf.fit, test)
mean((test$Count - rf.pred)^2)
```


```
## [1] 51281.56
```

```
importance(rf.fit)
```

```
##           %IncMSE IncNodePurity
## Hour          46.224894    456956560
## Temperature  19.153482    381116213
## Humidity     18.937996    180176308
## Visibility   12.229920     56458951
## Dew          9.259663    150428487
## Wind         7.636612     36426694
## Snowfall     4.064761     3198257
## Solar        12.975940    110319210
## Rainfall     11.555775     66623599
## Holiday      5.946961     3995529
## Month        9.036178    138860325
## Day         13.553250     42258375
## Functioning  34.479912    155387398
```

```
#Examine how random forest performs with different values of the tuning parameter mtry and see which value performs best
oob.err = double(10)
test.err = double(10)
for(mtry in 1:10)
{
  rf.count = randomForest(Count~Hour+Temperature+Humidity+Dew+Wind+Snowfall+Solar+Rainfall+Holiday+Month+Day+Functioning)
  oob.err[mtry]=rf.count$mse[50]
  rf.pred = predict(rf.count, test)
  test.err[mtry] = mean((test$Count - rf.pred)^2)
}
```

```
#Graph the output of the Out-of-Bag error and testing error for each value of mtry
matplot(1:mtry, cbind(test.err, oob.err), pch=19, col=c("green","blue"), type="b", xlab="mtry", ylab="MSE")
```



```
test.err[7]
```

```
## [1] 49531.73
```

Baseline Test

```
sample_mean = mean(out.data$Count)
mean((test$Count - sample_mean)^2)
```

```
## [1] 378820.8
```

Bagged Estimates

```
#For each of the models selected create new testing data to see how the model performs on average
bag_lin_mse = double(100)
for(i in 1:100)
{
  set.seed(i)
  sampleSize = floor(0.75 * nrow(data))
  split = sample(seq_len(nrow(data)), size = sampleSize)
  new_train = data[split, ]
  new_test = data[-split, ]
  lin.fit = lm(Count~Hour+Temperature+Humidity+Visibility+Solar+Rainfall+Holiday+Month+Day+Functioning,
  lin.pred = predict(lin.fit, new_test, type = 'response')
  bag_lin_mse[i] = mean((lin.pred - new_test$Count)^2)
}
avg = mean(bag_lin_mse)
avg
```

```
## [1] 149123.5
```

```

bag_lasso_mse = double(100)
for(i in 1:100)
{
  set.seed(i)
  sampleSize = floor(0.75 * nrow(data))
  split = sample(seq_len(nrow(data)), size = sampleSize)
  new_train = data[split, ]
  new_test = data[-split, ]
  train.x = model.matrix(Count~Hour+Temperature+Humidity+Visibility+Solar+Rainfall+Holiday+Month+Day+Fun
  train.apps = train$Count
  test.x = model.matrix(Count~Hour+Temperature+Humidity+Visibility+Solar+Rainfall+Holiday+Month+Day+Fun
  test.apps = new_test$Count
  lasso.fit = cv.glmnet(train.x, train.apps, alpha = 1)
  lasso.lambda = lasso.fit$lambda.min
  lasso.pred = predict(lasso.fit, s = lasso.lambda, newx = test.x)
  bag_lasso_mse[i] = mean((lasso.pred - test.apps)^2)
}
avg = mean(bag_lasso_mse)
avg

```

```
## [1] 149350.5
```

```

bag_rf_mse = double(100)
for(i in 1:100)
{
  set.seed(i)
  sampleSize = floor(0.75 * nrow(data))
  split = sample(seq_len(nrow(data)), size = sampleSize)
  new_train = data[split, ]
  new_test = data[-split, ]
  rf.fit = randomForest(Count~Hour+Temperature+Humidity+Visibility+Dew+Wind+Snowfall+Solar+Rainfall+Hol
  rf.pred = predict(rf.fit, new_test)
  bag_rf_mse[i] = mean((new_test$Count - rf.pred)^2)
}
avg = mean(bag_rf_mse)
avg

```

```
## [1] 19799.1
```

```

bag_gam_mse = double(100)
for(i in 1:100)
{
  set.seed(i)
  sampleSize = floor(0.75 * nrow(data))
  split = sample(seq_len(nrow(data)), size = sampleSize)
  new_train = data[split, ]
  new_test = data[-split, ]
  gam.mod = gam(Count~ s(Hour,5)+s(Temperature,5)+s(Humidity,4)+s(Visibility,4)+s(Solar,4)+s(Rainfall,5)
  preds = predict(gam.mod, new_test)
  bag_gam_mse[i] = mean((new_test$Count - preds)^2)
}
avg = mean(bag_gam_mse)
avg

```

```
## [1] 108874.7
```

Test Predictions

```
#Read and convert the data into the proper variable format  
test_data = read_csv("~/Documents/STAT1361/test.csv")
```

```
## Rows: 2208 Columns: 14
```

```
## -- Column specification -----  
## Delimiter: ","  
## chr (4): Date, Seasons, Holiday, Functioning  
## dbl (10): Hour, Temperature, Humidity, Wind, Visibility, Dew, Solar, Rainfal...
```

```
##  
## i Use 'spec()' to retrieve the full column specification for this data.  
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
test_data = test_data %>%  
  separate(Date, sep="/", into = c("Day", "Month", "Year"))
```

```
Count = predict(rf.fit, test_data)  
ID = test_data$ID  
student_id = rep(4293570, length(Count))  
test.pred = data.frame(ID, Count, student_id)
```